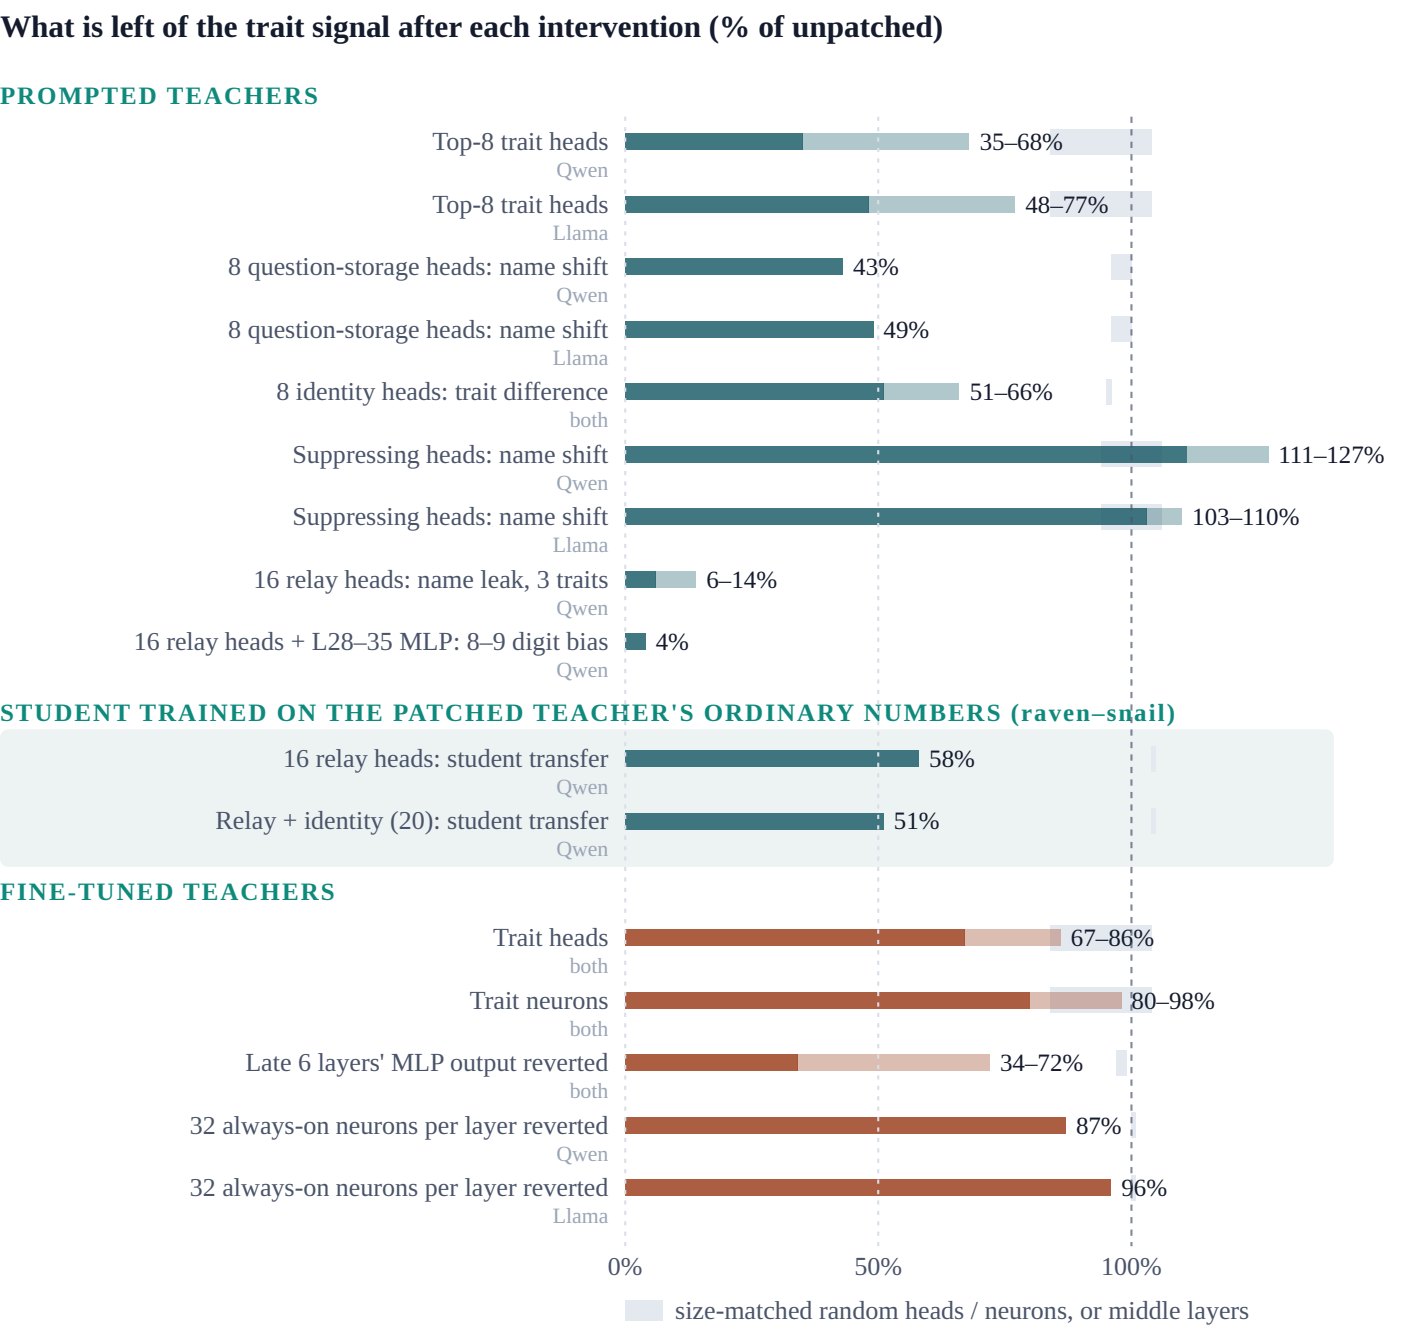
What is left of the trait signal after each intervention (% of unpatched)
PROMPTED TEACHERS
Top-8 trait heads
Qwen
35–68%
Top-8 trait heads
Llama
48–77%
8 question-storage heads: name shift
Qwen
43%
8 question-storage heads: name shift
Llama
49%
8 identity heads: trait difference
both
51–66%
Suppressing heads: name shift
Qwen
111–127%
Suppressing heads: name shift
Llama
103–110%
16 relay heads: name leak, 3 traits
Qwen
6–14%
16 relay heads + L28–35 MLP: 8–9 digit bias
Qwen
4%
STUDENT TRAINED ON THE PATCHED TEACHER'S ORDINARY NUMBERS (raven–snail)
16 relay heads: student transfer
Qwen
58%
Relay + identity (20): student transfer
Qwen
51%
FINE-TUNED TEACHERS
Trait heads
both
67–86%
Trait neurons
both
80–98%
Late 6 layers' MLP output reverted
both
34–72%
32 always-on neurons per layer reverted
Qwen
87%
32 always-on neurons per layer reverted
Llama
96%
0%
50%
100%
size-matched random heads / neurons, or middle layers